

# Distributed Systems

**Indirect  
communication**

# Contents

- Indirect vs. direct communication
- Types
- Event-based systems
  - Subscription models
  - IceStorm + MQTT
- Message queues systems
  - RabbitMQ

## Tip

All listings in this topic are available for download at:  
<https://github.com/UCLM-ESI/ssdd.examples>

# Indirect vs. direct communication

**Direct** communication is point-to-point:

- Participants must exist **at same time**.
  - Connection or session establishment
- Each participant require a way to know the address of other.
- Very inefficient when there are many participants.

Request-Reply, RPC/RMI protocols are direct communication.

# Indirect vs. direct communication

## Indirect communication:

- Communication implies an **intermediary**
- Sender/receiver(s) are decoupled in **time** and **space**:
  - In **space**: sender don't require to know the receiver identities or address, and vice-versa
    - sender/receiver(s) may be transparently replaced, migrated, replicated, etc.
  - In **time**: sender and receiver(s) may have different (even not overlapped) livetimes.
    - sender/receiver(s) don't require to coexist.

# Indirect vs. direct communication

	<i>Time-coupled</i>	<i>Time-uncoupled</i>
<i>Space coupling</i>	<p><i>Properties:</i> Communication directed towards a given receiver or receivers; receiver(s) must exist at that moment in time</p> <p><i>Examples:</i> Message passing, remote invocation (see Chapters 4 and 5)</p>	<p><i>Properties:</i> Communication directed towards a given receiver or receivers; sender(s) and receiver(s) can have independent lifetimes</p>
<i>Space uncoupling</i>	<p><i>Properties:</i> Sender does not need to know the identity of the receiver(s); receiver(s) must exist at that moment in time</p> <p><i>Examples:</i> IP multicast (see Chapter 4)</p>	<p><i>Properties:</i> Sender does not need to know the identity of the receiver(s); sender(s) and receiver(s) can have independent lifetimes</p> <p><i>Examples:</i> Most indirect communication paradigms covered in this chapter</p>

Instructor's Guide for Coulouris, Dollimore, Kindberg and Blair, Distributed Systems: Concepts and Design Edn. 5  
 © Pearson Education 2012

# Suitable for...

- scenarios where users **connect** and **disconnect** very often.
- scenarios with a **large number** of users
- **event dissemination** where recipients are unknown and change often
  - Mobile environments, messaging services, notification
  - Example: RSS, IoT, etc.

# Disadvantages

- **Overloading** due to the introduction of a level of indirection
  - Reliable delivery of messages, order, etc.
- Increased **management complexity** due to decoupling between transmitters/receivers.
- **Difficulty in achieving** end-to-end properties
  - Real time, security, etc.

# Indirect communication

Communication between the sender and the receiver through an **intermediary**:

## Group communication

- Senders send messages to a "group". They don't know the recipient identities.

## Distributed event-based systems

- Disseminate events to multiple recipients through an intermediary (the event broker).

## Message queues

- Senders put messages to a distributed queue, and recipients extract them from the queue.


## Distributed shared memory

- Simulate centralized shared memory.



# Group communication

- Central concept: GROUP
- Processes may become group members with **join()**, then may **leave()**.
- To send a message to a group:  
**group.send(message)**
- Very complex when reliability is required



# Distributed event-based systems

Distributed Event-based systems implement one-to-many communications

- Heterogeneity & Asynchronism
- More decoupled and reactive style than RMI and/or RPC
- **Public-subscribe systems**

Roles:

- **Publisher** publishes structured events to an event service
- **Subscriber** expresses interest in particular events through subscriptions to an event service
- **Event service or broker** receives events from publishers and delivers events to subscribers according to its interests.

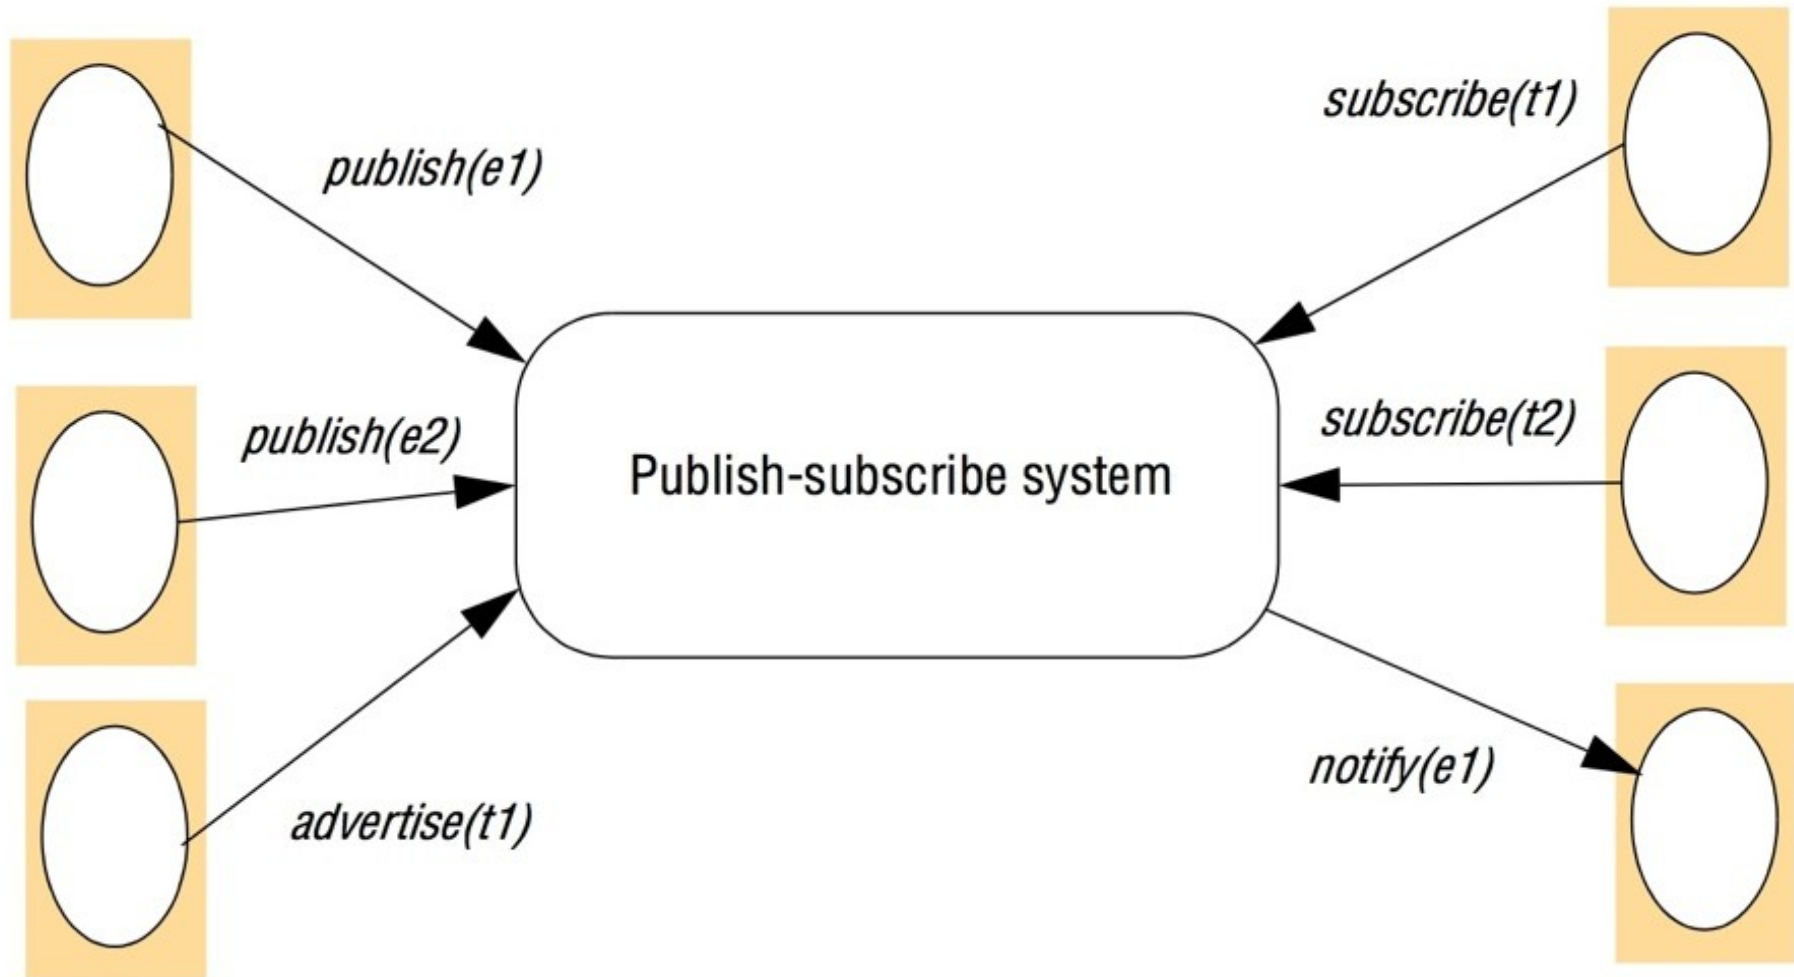
# Programming model

- $e \rightarrow$  Event
  - any structured piece of information
- $f \rightarrow$  Filter
  - A pattern to filter the type of events
  - Expressed in specific format (e.g. regular expressions) or language (e.g. SQL-like languages)
- Basic operations :
  - `advertise(f)` / `unadvertise(f)`
  - `subscribe(f)` / `unsubscribe(f)`
  - `publish(e)` / `notify(e)`

# Publish-subscribe paradigm

Publishers


Subscribers



# Subscription models


**Channel-based:** *publishers* publish events to **named channels** and *subscribers* then subscribe to one of these named channels to receive all events sent to the channel.

- Ex: ZeroC IceStorm



**Topic-based (or subject based):** Similar to channel-based but a topic (may be hierarchical) is explicitly defined as a **field**.


- Ex: MQTT




# Subscription models

**Content-based:** generalization of topic-based where the filter is a query defined in terms of compositions of constraints over the values of event attributes.

- Ex: DDS




**Type-based:** subscriptions are defined in terms of types of events and matching is defined in terms of types or subtypes of the given filter.



# Centralized vs. Distributed

## Centralized

- The *broker* is implemented in a single node
- **Pros:** straightforward to implement
- **Cons:** lack of resilience and scalability



## Distributed

- A *network of brokers* cooperate for event distribution

## Description: A minimal MQTT working example

- `examples/mqtt`

Assure broker is running:

```
$ sudo service mosquitto restart
```

See and play with:

- `subscriber.py`
- `publisher-humidity.py`
- `publisher-temperature.py`

Run publisher(s)

```
$ python3 publisher-temperature.py &  
python3 publisher-humidity.py
```

```
$ python3 subscriber.py  
topic: temperature/X002, msg:  
{  
  'identifier': 'X002',  
  'value': 36, 'unit':  
  'Celsius',  
  'timestamp': 1668347911.566947  
} 36
```

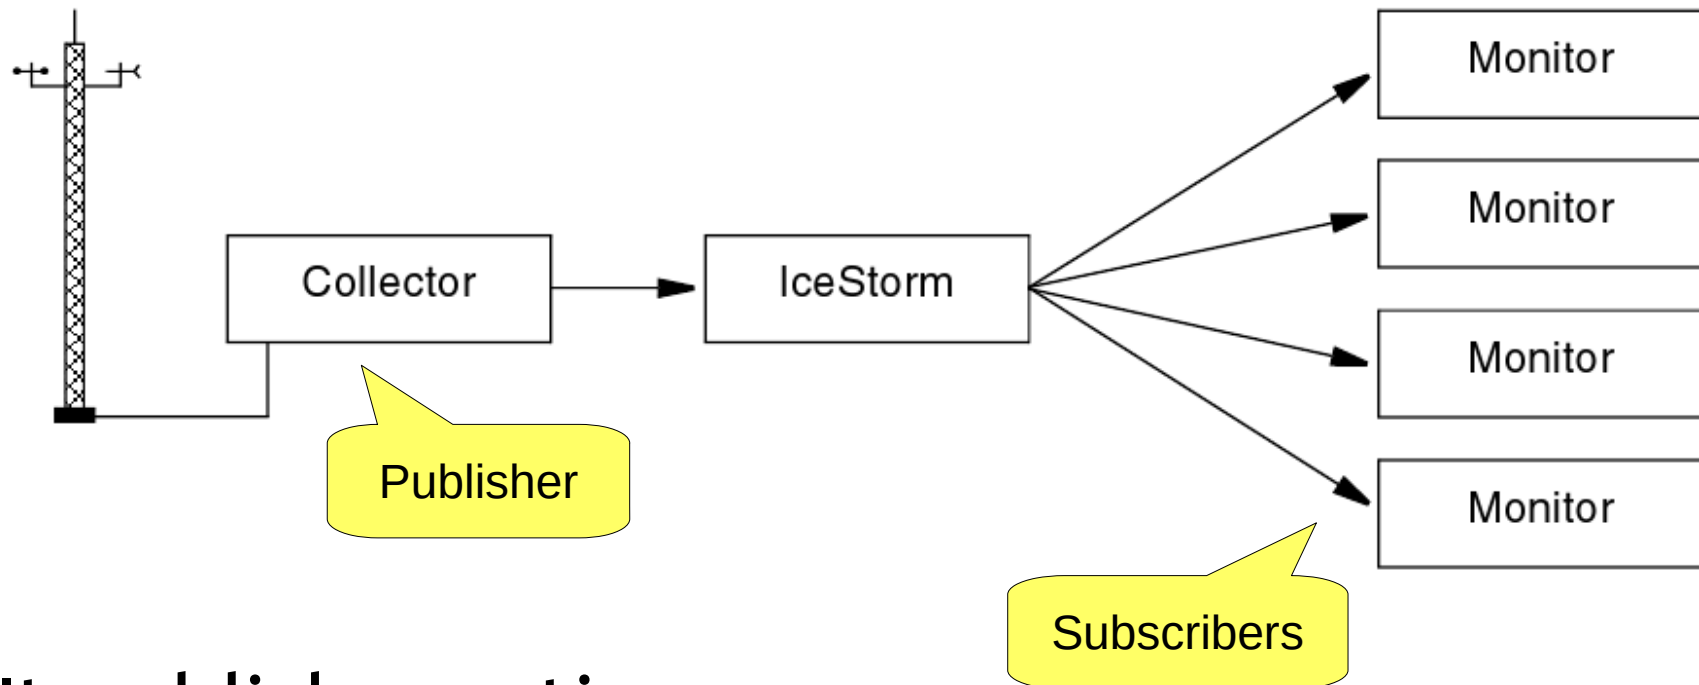
Run subscriber(s)  
(try several)



It is the ZeroC Ice event service

- Invocation oriented (vs. data oriented)
- Channel based (named “topics”)
- Centralized broker (but replicated: HA-IceStorm)
- Basic filter features

# IceStorm




- It publish one time:
  - Broker creates copies
- Publish = invoke oneway method
  - Only provides push

# IceStorm

## Basic features

- Decouples production and consumption
- Allows to use a different transport protocol for any publisher or subscriber.
- Federation



Description: A minimal IceStorm working example

- `hello.ice/icestorm`

Choose branch:

```
$ git checkout branches/ice37
```

Start broker:

```
$ make start  
icebox --Ice.Config=icebox.config &
```

See and play with:

- `subscriber.py`
- `publisher.py`

```
$ make run-subscriber  
./subscriber.py --Ice.Config=subscriber.config  
Using IceStorm in: 'IceStorm.TopicManager.Proxy'  
Waiting events... 'CB1EF114-2937-4F8A-BADD-882EAAE203CC ...  
Event received: Hello World 0!
```

Run subscriber(s)  
(try several)

Run publisher(s)

```
$ make run-publisher  
./publisher.py --Ice.Config=publisher.config  
Using IceStorm in: 'IceStorm.TopicManager.Proxy'  
publishing 10 'Hello World' events
```

# Message queues systems

- Distributed message queues decoupled **one-to-one** communication.
  - A **sender** put a message into an **explicit** queue.
  - A (single) **recipient** remove the message from queue.
- Commercial middlewares:
  - WebSphere MQ, Oracle AQ
  - JMS (Java Message Service)/Open MQ


# Message queues

- **Queues** are usually **FIFO (First In First Out)**, but:
  - Some implementations provide priority or selection criterion.
- Messages consist of:
  - Destination queue identifier
  - Meta data
  - Payload
- Features:
  - Validity (message will be received)
  - Integrity (just one time)
  - Persistent

# Programming model

Three styles for receiving messages:

- blocking receive (**pull**)
- non-blocking receive (**polling**)
- notify (**push**)



**RabbitMQ** implements the standard **AMQP** (**A**dvanced **M**essage **Q**ueuing **P**rotocol)

Concepts:

- **Publishers:** send messages to an exchange
- **Exchanges:** implement routing to send messages to consumers.
- **Consumers:** declare queue and bind it to a exchange to receive messages.
- **Queues:** Temporarily store messages from publishers
- **Routing:** Apply routing keys to perform message matching:  
*direct, fan-out* and *topic* exchanges.


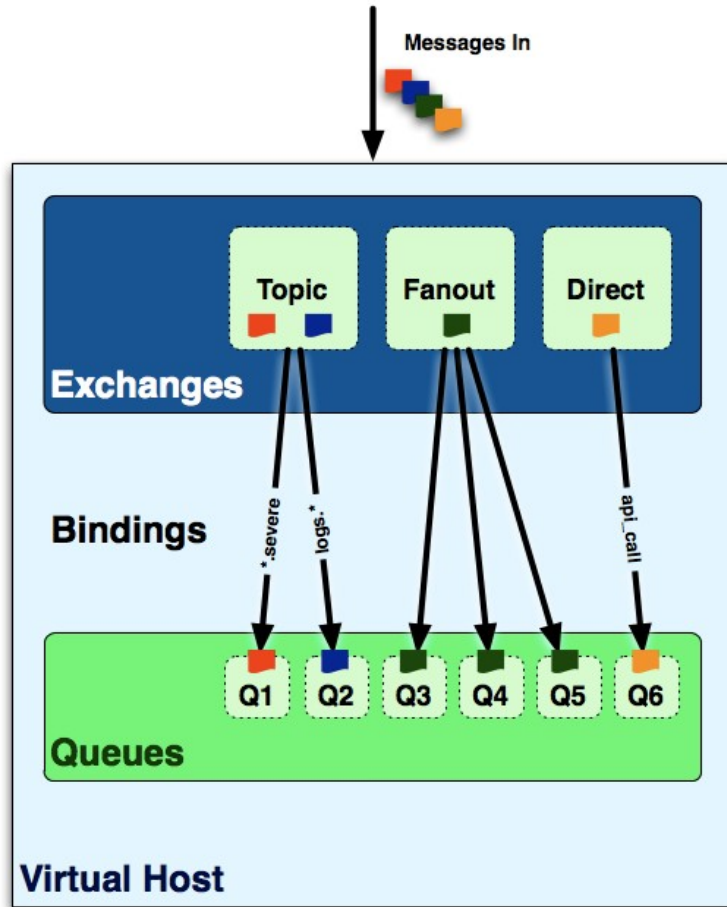
It supports several messaging protocols:

- AMQP, MQTT, HTTP...



# RabbitMQ

## Message Queue example



Description: minimal message queue working example

- `[examples:rabbit-hello]`



consumer.py

```
import pika

def callback(ch, method, properties, body):
    print " [x] Received %r" % (body,)

localhost = pika.ConnectionParameters(host='localhost')
connection = pika.BlockingConnection()
channel = connection.channel()
channel.queue_declare(queue='hello')
channel.basic_consume(callback, queue='hello', no_ack=True)

print('[*] Waiting for messages. To exit press CTRL+C')
channel.start_consuming()
```

Description: minimal message queue working example

- `[examples:rabbit-hello]`



producer.py

```
#!/usr/bin/env python
import pika


localhost = pika.ConnectionParameters(host='localhost')
connection = pika.BlockingConnection()
channel = connection.channel()

channel.basic_publish(exchange='',
                      routing_key='hello',
                      body='Hello World!')


print " [x] Sent 'Hello World!'"
connection.close()
```

# RabbitMQ


## flexible binding




work queues



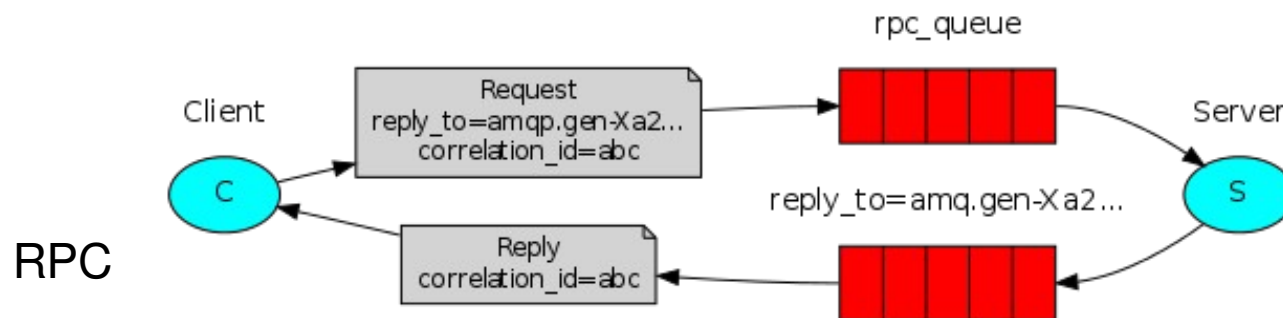
publish-subscribe



routing



topics



RPC

# References

G. Coulouris, *Distributed Systems: Concepts and Design*, Addison Wesley 2011, Fifth edition

- Section 6.1 – Introduction
- Section 6.3 – Publish-subscribe systems
- Section 6.4 – Message queues